

Final Project Submission

Please fill out:

- Student name: Leshmi Jayakumar
- Student pace: Part time
- Scheduled project review date/time: 12/02/2023
- Instructor name: Hardik Idnani

KING COUNTRY HOUSE DATA ANALYSIS

****Column Names and descriptions for Kings County Data Set**

- **date****Date** - house was sold
- **price****Price** - is prediction target
- **bedrooms****Number** - of Bedrooms/House
- **bathrooms****Number** - of bathrooms/bedrooms
- **sqft_livingsquare** - footage of the home
- **sqft_lotsquare** - footage of the lot
- **floors****Total** - floors (levels) in house
- **waterfront** - House which has a view to a waterfront
- **view** - Has been viewed
- **condition** - How good the condition is (Overall)
- **grade** - overall grade given to the housing unit, based on King County grading system
- **sqft_above** - square footage of house apart from basement
- **sqft_basement** - square footage of the basement
- **yr_built** - Built Year
- **yr_renovated** - Year when house was renovated
- **zipcode** - zip
- **lat** - Latitude coordinate
- **long** - Longitude coordinate
- **sqft_living15** - The square footage of interior housing living space for the nearest 15 neighbors
- **sqft_lot15** - The square footage of the land lots of the nearest 15 neighbors

Goal

The goal of this project is to create a relatively accurate prediction model for the prices that future houses sell for. We will explore how different factors affect the pricing the homes, given the data, and compile the most important feature.

In [1]:



```
# First, let's import the libraries we are going to use
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline

# Ignores warnings
import warnings
warnings.filterwarnings('ignore')

# Used for working with the z-score
from scipy import stats
from scipy import interpolate
from itertools import combinations
#For OLS model
import statsmodels.api as sm
import statsmodels.formula.api as smf
import scipy.stats as stats
import statsmodels.stats.api as sms
from statsmodels.formula.api import ols
import pylab

# Used for Linear Regression model and Cross Validation model
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import mean_squared_error
from sklearn.metrics import mean_squared_error, make_scorer
from sklearn.model_selection import cross_val_score
from sklearn import metrics

# Used when showing normalization of data
from statsmodels.stats.diagnostic import normal_ad

# Used when creating bar plots and using median instead of mean
from numpy import median
```

In [2]:



```
#Let's import our King's County housing data and take a look
data_k=pd.read_csv('data/kc_house_data.csv')
```

In [3]:

```
data_k.head()
```

Out[3]:

	id	date	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterf
0	7129300520	10/13/2014	221900.0	3	1.00	1180	5650	1.0	
1	6414100192	12/9/2014	538000.0	3	2.25	2570	7242	2.0	
2	5631500400	2/25/2015	180000.0	2	1.00	770	10000	1.0	
3	2487200875	12/9/2014	604000.0	4	3.00	1960	5000	1.0	
4	1954400510	2/18/2015	510000.0	3	2.00	1680	8080	1.0	

5 rows × 21 columns

In [4]:

```
#Data information
data_k.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21597 non-null  int64
1   date                  21597 non-null  object
2   price                 21597 non-null  float64
3   bedrooms              21597 non-null  int64
4   bathrooms             21597 non-null  float64
5   sqft_living           21597 non-null  int64
6   sqft_lot              21597 non-null  int64
7   floors                21597 non-null  float64
8   waterfront            19221 non-null  float64
9   view                  21534 non-null  float64
10  condition             21597 non-null  int64
11  grade                 21597 non-null  int64
12  sqft_above            21597 non-null  int64
13  sqft_basement         21597 non-null  object
14  yr_built              21597 non-null  int64
15  yr_renovated          17755 non-null  float64
16  zipcode               21597 non-null  int64
17  lat                   21597 non-null  float64
18  long                  21597 non-null  float64
19  sqft_living15         21597 non-null  int64
20  sqft_lot15            21597 non-null  int64
dtypes: float64(8), int64(11), object(2)
memory usage: 3.5+ MB
```

In [5]:

```
#Convert date to date time
data_k['date'] = pd.to_datetime(data_k['date'], format='%m/%d/%Y')
#Convert sqft_basement to float
data_k['sqft_basement'] = pd.to_numeric(data_k['sqft_basement'], errors="coerce")
#using errors='coerce' because sqft_basement contains '?' string values. These values will
```

In [6]:

```
#after converting the dtype of sqft_basement and Date
data_k.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   id                    21597 non-null  int64
1   date                 21597 non-null  datetime64[ns]
2   price               21597 non-null  float64
3   bedrooms            21597 non-null  int64
4   bathrooms           21597 non-null  float64
5   sqft_living         21597 non-null  int64
6   sqft_lot            21597 non-null  int64
7   floors              21597 non-null  float64
8   waterfront          19221 non-null  float64
9   view                21534 non-null  float64
10  condition            21597 non-null  int64
11  grade               21597 non-null  int64
12  sqft_above          21597 non-null  int64
13  sqft_basement       21143 non-null  float64
14  yr_built            21597 non-null  int64
15  yr_renovated        17755 non-null  float64
16  zipcode             21597 non-null  int64
17  lat                 21597 non-null  float64
18  long                21597 non-null  float64
19  sqft_living15       21597 non-null  int64
20  sqft_lot15          21597 non-null  int64
dtypes: datetime64[ns](1), float64(9), int64(11)
memory usage: 3.5 MB
```

In [7]:



```
# Handle the missing data.  
  
print(data_k.isnull().sum())  
print(data_k.shape)
```

```
id                0  
date              0  
price             0  
bedrooms          0  
bathrooms         0  
sqft_living       0  
sqft_lot          0  
floors            0  
waterfront       2376  
view              63  
condition         0  
grade             0  
sqft_above        0  
sqft_basement     454  
yr_built          0  
yr_renovated      3842  
zipcode           0  
lat               0  
long              0  
sqft_living15     0  
sqft_lot15        0  
dtype: int64  
(21597, 21)
```

In [8]:



```
#For sqft_basement, waterfront and view, replace missing values with 0  
data_k['sqft_basement'].fillna(0, inplace=True)  
data_k['waterfront'].fillna(0, inplace=True)  
data_k['view'].fillna(0, inplace=True)  
#For yr_renovated, set yr_renovated to yr_built  
data_k['yr_renovated'].fillna(data_k[data_k['yr_renovated'].isna()]['yr_built'], inplace=True)
```

In [9]:



```
print(data_k.isnull().sum())
```

```
id          0
date        0
price       0
bedrooms    0
bathrooms   0
sqft_living  0
sqft_lot    0
floors       0
waterfront  0
view         0
condition    0
grade        0
sqft_above  0
sqft_basement 0
yr_built     0
yr_renovated 0
zipcode      0
lat          0
long         0
sqft_living15 0
sqft_lot15   0
dtype: int64
```

Questions

Will the price of the house positively correlated with :

- * a)No:of bedrooms
- * b)Sqft_living
- * c)Condition
- * d)Grade
- * e)Zipcode

Baseline

In [10]:



```
#Let's describe our data now that all our values have been cleaned up.  
data_k.describe()
```

Out[10]:

	id	price	bedrooms	bathrooms	sqft_living	sqft_lot
count	2.159700e+04	2.159700e+04	21597.000000	21597.000000	21597.000000	2.159700e+04
mean	4.580474e+09	5.402966e+05	3.373200	2.115826	2080.321850	1.509941e+04
std	2.876736e+09	3.673681e+05	0.926299	0.768984	918.106125	4.141264e+04
min	1.000102e+06	7.800000e+04	1.000000	0.500000	370.000000	5.200000e+02
25%	2.123049e+09	3.220000e+05	3.000000	1.750000	1430.000000	5.040000e+03
50%	3.904930e+09	4.500000e+05	3.000000	2.250000	1910.000000	7.618000e+03
75%	7.308900e+09	6.450000e+05	4.000000	2.500000	2550.000000	1.068500e+04
max	9.900000e+09	7.700000e+06	33.000000	8.000000	13540.000000	1.651359e+06

Looking at the information above we can see the following:

1) Our target variable, Price, has a mean value of roughly 540K. The lowest price is 78K and the highest price is over 7 million. We can also see that 50% of our data falls below the 500K price. 2) We can also see some outliers such as a home with 33 bedrooms and 8 bathrooms. 3) We also notice that some of our data is categorical, such as zipcode and waterfront.

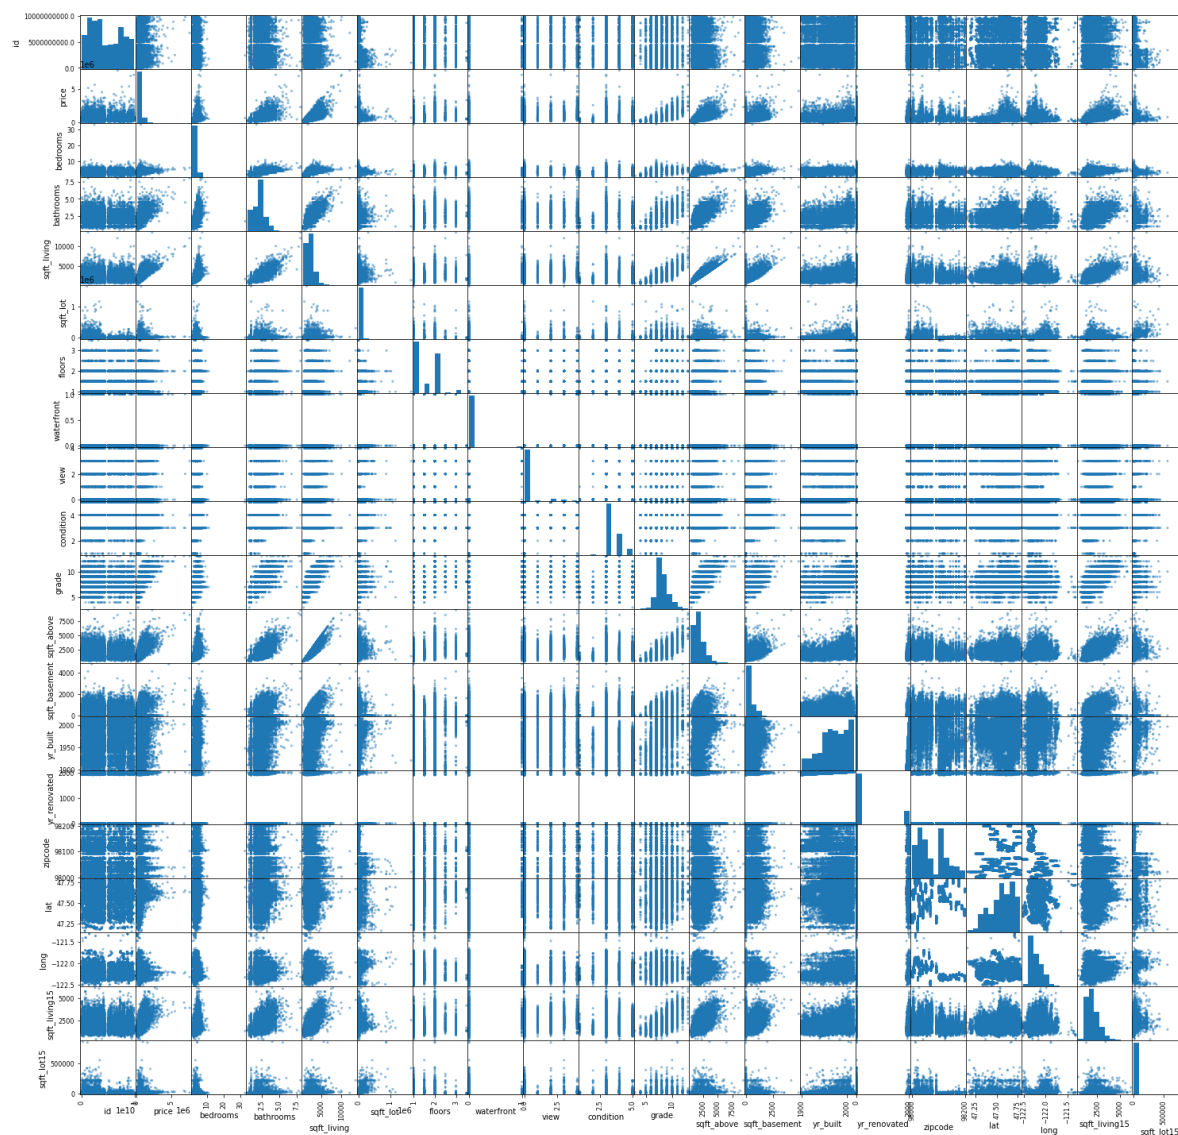
Now, let's run some graphics in order to better look at our data

In [11]:



```
#Scatter plot of the data
```

```
pd.plotting.scatter_matrix(data_k, figsize=(25,25));
```



Looking at our scatter matrix above, we can see interesting relationships in our data, but because the data is difficult to actually look at.

In [12]:

```
#checking the mean and standard deviation
price_mean = data_k['price'].mean()
price_std = data_k['price'].std()
price_mean, price_std
```

Out[12]:

```
(540296.5735055795, 367368.1401013945)
```

In [13]:

```
#checking the average price of the house according to the year built
avg_price_yr=data_k.groupby(['yr_built']).price.mean()
avg_price_yr.sort_index()
```

Out[13]:

```
yr_built
1900    581536.632184
1901    557108.344828
1902    673192.592593
1903    480958.195652
1904    583867.755556
...
2011    544648.384615
2012    527436.982353
2013    678599.582090
2014    683792.685152
2015    759970.947368
Name: price, Length: 116, dtype: float64
```


In [14]:

```
#Data visulisation according to the price and year built
T=avg_price_yr.index
X=np.linspace(T.min(),T.max(),10000)
Y=avg_price_yr.values
spl=interpolate.make_interp_spline(T,Y,3)
test=spl(X)

fig,ax=plt.subplots(figsize=(30,15))
ax.plot(X,test)
# Interpolate to make more smooth(Interpolation in Python is a technique used to estimate u
```

Out[14]:

[<matplotlib.lines.Line2D at 0x2a49bb75880>]



In [299]:

Answer:This figure shows that the year built doesnot affect the price and it also showsther

File "<ipython-input-299-622f73eb687e>", line 1

Answer:This figure shows that the year built doesnot affect the price


SyntaxError: invalid syntax

In [15]:

```
fig, ax = plt.subplots(figsize=(12,12))
ax = sns.distplot(data_k['price'], label='Price Density')
plt.axvline(price_mean, color='red', label='Mean Price')
plt.legend()
```

Out[15]:

<matplotlib.legend.Legend at 0x2a4a589c580>



In [300]:



```
data_k.corr()
```

Out[300]:

id	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
1.000000	-0.016772	0.001150	0.005162	-0.012241	-0.131911	0.018608	-0.003599	0.01177
-0.016772	1.000000	0.308787	0.525906	0.701917	0.089876	0.256804	0.264306	0.39349
0.001150	0.308787	1.000000	0.514508	0.578212	0.032471	0.177944	-0.002127	0.07835
0.005162	0.525906	0.514508	1.000000	0.755758	0.088373	0.502582	0.063629	0.18601
-0.012241	0.701917	0.578212	0.755758	1.000000	0.173453	0.353953	0.104637	0.28171
-0.131911	0.089876	0.032471	0.088373	0.173453	1.000000	-0.004814	0.021459	0.07505
0.018608	0.256804	0.177944	0.502582	0.353953	-0.004814	1.000000	0.020797	0.02841
-0.003599	0.264306	-0.002127	0.063629	0.104637	0.021459	0.020797	1.000000	0.38054
0.011772	0.393497	0.078354	0.186016	0.281715	0.075054	0.028414	0.380543	1.00000
-0.023803	0.036056	0.026496	-0.126479	-0.059445	-0.008830	-0.264075	0.016648	0.04562
0.008188	0.667951	0.356563	0.665838	0.762779	0.114731	0.458794	0.082818	0.24908
-0.010799	0.605368	0.479386	0.686668	0.876448	0.184139	0.523989	0.071778	0.16601
-0.004359	0.321108	0.297229	0.278485	0.428660	0.015031	-0.241866	0.083050	0.27062
0.021617	0.053953	0.155670	0.507173	0.318152	0.052946	0.489193	-0.024487	-0.05445
-0.007119	0.051588	-0.002292	0.012400	0.016479	-0.006170	-0.001467	0.031756	0.04483
-0.008211	-0.053402	-0.154092	-0.204786	-0.199802	-0.129586	-0.059541	0.028923	0.08505
-0.001798	0.306692	-0.009951	0.024280	0.052155	-0.085514	0.049239	-0.012157	0.00632
0.020672	0.022036	0.132054	0.224903	0.241214	0.230227	0.125943	-0.037628	-0.07770
-0.002701	0.585241	0.393406	0.569884	0.756402	0.144763	0.280102	0.083823	0.27892
-0.138557	0.082845	0.030690	0.088303	0.184342	0.718204	-0.010722	0.030658	0.07308




In [16]:

```
#checking the correlation using data visualization
plt.figure(figsize=(15,10))
sns.heatmap(data_k.corr().round(2),annot=True, cmap=('YlGnBu'))
```

Out[16]:

<AxesSubplot:>



But, even our heatmap is difficult to read. We are going to use a pairwise correlation in order to determine which values are highly correlated with our target (price) by using any value greater than 0.70.


Checking outliers

In [17]:

```
plt.figure(figsize=(12,8))  
sns.boxplot(x='bedrooms',y='price',data=data_k)
```

Out[17]:

<AxesSubplot:xlabel='bedrooms', ylabel='price'>




In [18]:

```
plt.figure(figsize=(12,8))
sns.boxplot(x='bathrooms',y='price',data=data_k)
```

Out[18]:

<AxesSubplot:xlabel='bathrooms', ylabel='price'>




In [19]:

```
sns.distplot(data_k.sqft_living)
```

Out[19]:

<AxesSubplot:xlabel='sqft_living', ylabel='Density'>




In [20]:

```
plt.figure(figsize=(12,8))
sns.boxplot(x='floors',y='price',data=data_k)
```

Out[20]:

<AxesSubplot:xlabel='floors', ylabel='price'>



In [21]:

```
plt.figure(figsize=(10,7))
sns.scatterplot(data_k['condition'], data_k['price'])
plt.title('House Condition and price', fontsize=15, fontname='silom')
```

Out[21]:

Text(0.5, 1.0, 'House Condition and price')

findfont: Font family ['silom'] not found. Falling back to DejaVu Sans.


The house condition is not affecting the price according to this figure

In [22]:

```
plt.figure(figsize=(10,7))
sns.scatterplot(data_k['condition'], data_k['grade'])
plt.title('House Condition and Grade', fontsize=15, fontname='silom')
```

Out[22]:

Text(0.5, 1.0, 'House Condition and Grade')



In []:

In this figure the condition of the house doesn't affect the grade

In [287]:

```
plt.figure(figsize=(12,8))
sns.boxplot(x='grade',y='price',data=data_k)
```

Out[287]:

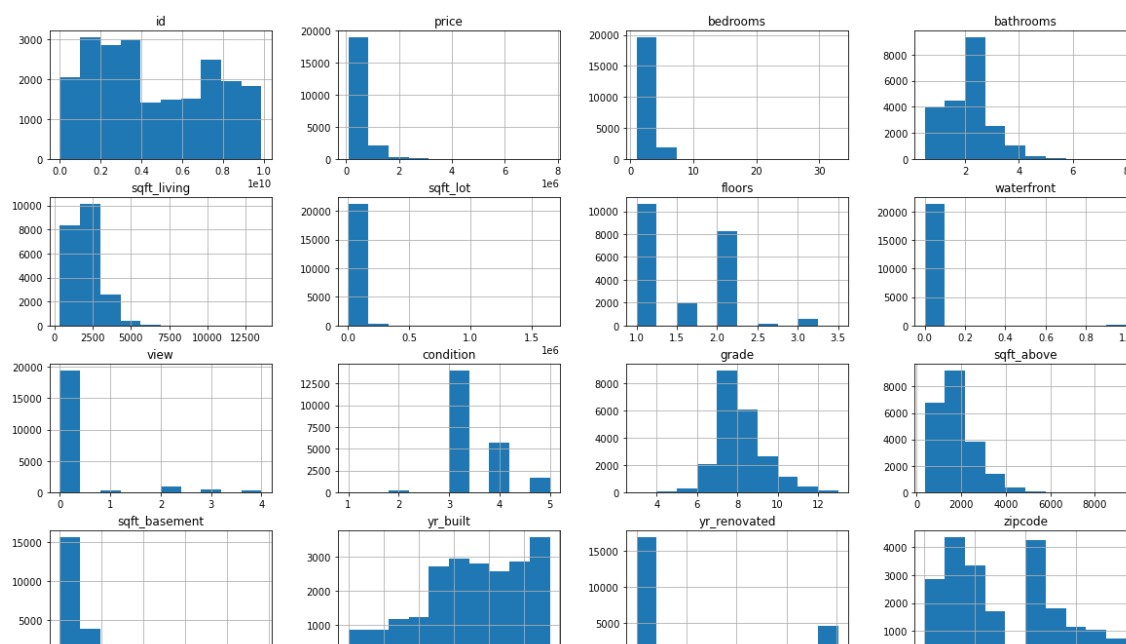
<AxesSubplot:xlabel='grade', ylabel='price'>



In this figure the house price increase with increase in grade

In [23]:

```
fig = plt.figure(figsize = (20,15))
ax = fig.gca()
data_k.hist(ax = ax);
```



In []:

We can see from the box plots and histograms above that many of our continuous data have out

In [303]:

```
outcome = 'price'
x_cols = data_k.drop('price',axis=1)
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=data_k).fit()
model.summary()
```

Out[303]:

OLS Regression Results

Dep. Variable:	price	R-squared:	0.707
Model:	OLS	Adj. R-squared:	0.702
Method:	Least Squares	F-statistic:	131.8
Date:	Sun, 12 Feb 2023	Prob (F-statistic):	0.00
Time:	22:07:35	Log-Likelihood:	-2.9412e+05
No. Observations:	21597	AIC:	5.890e+05
Df Residuals:	21207	BIC:	5.921e+05
Df Model:	389		
Covariance Type:	nonrobust		

In [307]:

```
new_data= data_k.drop(['id','view','date'], axis=1)
new_data
```

Out[307]:

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	condition	grade
0	221900.0	3	1.00	1180	5650	1.0	0.0	3	7
1	538000.0	3	2.25	2570	7242	2.0	0.0	3	7
2	180000.0	2	1.00	770	10000	1.0	0.0	3	6
3	604000.0	4	3.00	1960	5000	1.0	0.0	5	7
4	510000.0	3	2.00	1680	8080	1.0	0.0	3	8
...
1592	360000.0	3	2.50	1530	1131	3.0	0.0	3	8
1593	400000.0	4	2.50	2310	5813	2.0	0.0	3	8
1594	402101.0	2	0.75	1020	1350	2.0	0.0	3	7
1595	400000.0	3	2.50	1600	2388	2.0	0.0	3	8
1596	325000.0	2	0.75	1020	1076	2.0	0.0	3	7

597 rows × 18 columns



In [309]:

```
new_data['waterfront'].fillna(0, inplace=True)
```

In [310]:

```
#Now, let's place our continuous and categorical data into a separate features
cont_data = new_data[['price', 'sqft_living', 'sqft_lot', 'sqft_basement', 'yr_built', 'lat
']]
cat_data = new_data[['bedrooms', 'bathrooms', 'floors', 'waterfront', 'condition', 'grade', 'z
```

Remove outliers

In [311]:

```
count = 0
bath_outliers = []
mean = np.mean(cat_data['bathrooms'])
max_distance = np.std(cat_data['bathrooms']) * 3

for idx, row in cat_data['bathrooms'].T.iteritems():
    if abs(row-mean) >= max_distance:
        count += 1
        cat_data.drop(idx, inplace=True)
count
```

Out[311]:

187

In [312]:

```
count = 0
bed_outliers = []
mean = np.mean(cat_data['bedrooms'])
max_distance = np.std(cat_data['bedrooms']) * 3

for idx, row in cat_data['bedrooms'].T.iteritems():
    if abs(row-mean) >= max_distance:
        count += 1
        cat_data.drop(idx, inplace=True)
count
```

Out[312]:

47

In [313]:

```
cont_data.shape
```

Out[313]:

(21597, 11)

In [314]:

```
z_cont=np.abs(stats.zscore(cont_data))
```

In [315]:

```
cont_data1 = cont_data[(z_cont < 3).all(axis=1)]
```

In [316]:

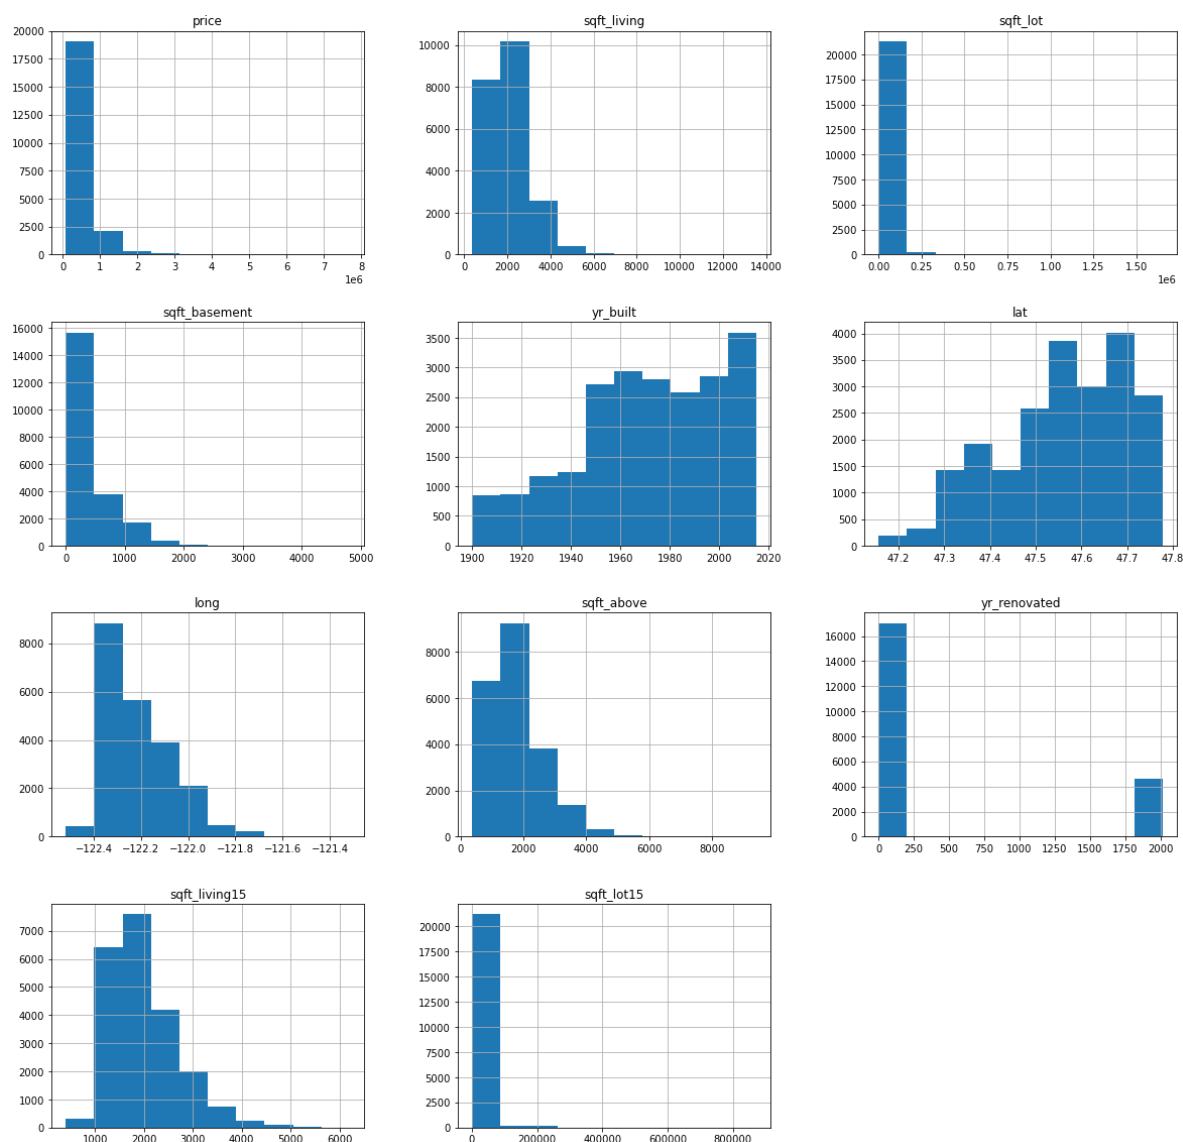
```
cont_data1.shape
```

Out[316]:

(20135, 11)

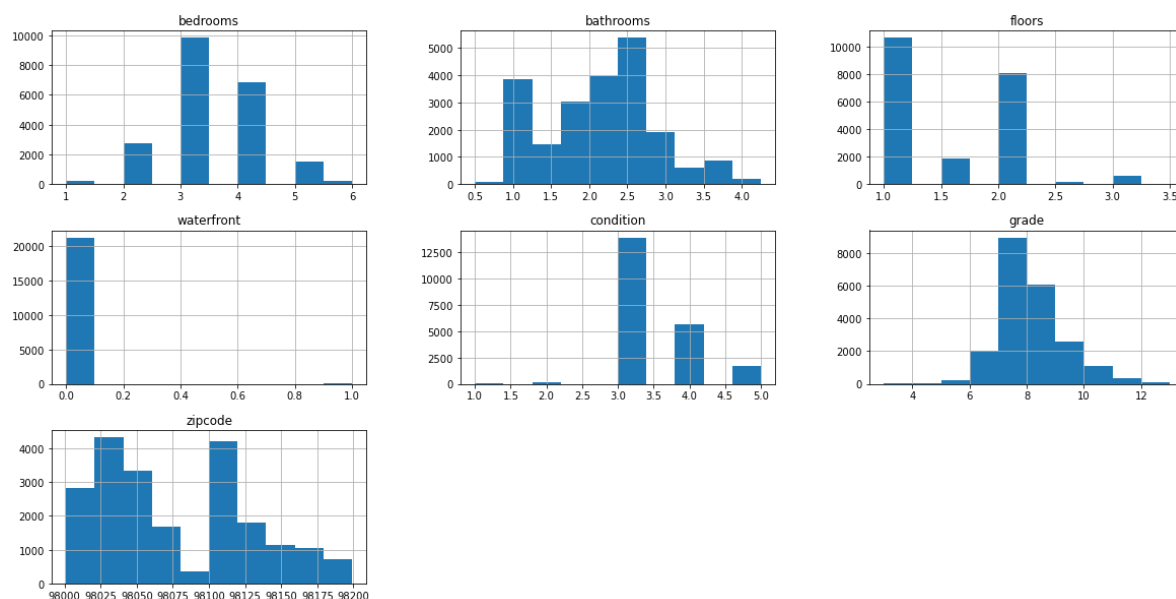
In [317]:

```
# Analysing data with histogram  
cont_data.hist(figsize=(20,20));
```



In [107]:

```
# Analysing data with histogram
cat_data.hist(figsize=(20,10));
```



In [289]:

```
plt.figure(figsize=(12,8))
sns.boxplot(x='sqft_living', y='price', data=cont_data)
```

Out[289]:

<AxesSubplot:xlabel='sqft_living', ylabel='price'>



In [318]:

```
#Let's make a continuous data feature and a feature for our target variable
target = cont_data['price']
cont_feat = cont_data.drop('price', axis=1)
```

In [319]:

```
cont_feat.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 21597 entries, 0 to 21596
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype  
---  -
0   sqft_living            21597 non-null  int64  
1   sqft_lot               21597 non-null  int64  
2   sqft_basement          21597 non-null  float64 
3   yr_built               21597 non-null  int64  
4   lat                   21597 non-null  float64 
5   long                  21597 non-null  float64 
6   sqft_above            21597 non-null  int64  
7   yr_renovated           21597 non-null  float64 
8   sqft_living15          21597 non-null  int64  
9   sqft_lot15            21597 non-null  int64  
dtypes: float64(4), int64(6)
memory usage: 1.6 MB
```

In [320]:

```
target
```

Out[320]:

```
0      221900.0
1      538000.0
2      180000.0
3      604000.0
4      510000.0
...
21592   360000.0
21593   400000.0
21594   402101.0
21595   400000.0
21596   325000.0
Name: price, Length: 21597, dtype: float64
```

In [321]:

```
dummy_1= pd.get_dummies(data=cat_data, columns=['bedrooms'],prefix='bed', drop_first=True)
dummy_2= pd.get_dummies(data=cat_data, columns=['condition'],prefix='condi', drop_first=True)
dummy_3= pd.get_dummies(data=cat_data, columns=['grade'],prefix='grd', drop_first=True)
dummy_4= pd.get_dummies(data=cat_data, columns=['zipcode'],prefix='zip_c', drop_first=True)
#dummy_5= pd.get_dummies(data=cat_data, columns=['waterfront'],prefix='wf', drop_first=True)
#dummy_6= pd.get_dummies(data=cat_data, columns=['floors'],prefix='flr', drop_first=True)
#dummy_7= pd.get_dummies(data=cat_data, columns=['bathrooms'],prefix='bath', drop_first=True)
```

In [322]:

```
#Let's create one feature containing all of our dummies
cat_dummies = pd.concat([dummy_1,dummy_2,dummy_3,dummy_4],axis=1)
cat_dummies
```

	bathrooms	floors	waterfront	condition	grade	zipcode	bed_2	bed_3	bed_4	bed_5	...	zip_c_981
0	1.00	1.0	0.0	3	7	98178	0	1	0	0	...	
1	2.25	2.0	0.0	3	7	98125	0	1	0	0	...	
2	1.00	1.0	0.0	3	6	98028	1	0	0	0	...	
3	3.00	1.0	0.0	5	7	98136	0	0	1	0	...	
4	2.00	1.0	0.0	3	8	98074	0	1	0	0	...	
...	
21592	2.50	3.0	0.0	3	8	98103	0	1	0	0	...	
21593	2.50	2.0	0.0	3	8	98146	0	0	1	0	...	
21594	0.75	2.0	0.0	3	7	98144	1	0	0	0	...	
21595	2.50	2.0	0.0	3	8	98027	0	1	0	0	...	
21596	0.75	2.0	0.0	3	7	98144	1	0	0	0	...	

In [323]:

```
d1= pd.concat([target,cont_feat,cat_dummies], axis=1)
d1.dropna(how='any',inplace=True)
```


In [324]:



```
outcome = 'price'
x_cols = d1.drop('price',axis=1)
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=d1).fit()
model.summary()
```

Out[324]:

OLS Regression Results

Dep. Variable:	price		R-squared:		0.822	
Model:	OLS		Adj. R-squared:		0.821	
Method:	Least Squares		F-statistic:		969.6	
Date:	Sun, 12 Feb 2023		Prob (F-statistic):		0.00	
Time:	22:17:30		Log-Likelihood:		-2.8333e+05	
No. Observations:	21363		AIC:		5.669e+05	
Df Residuals:	21261		BIC:		5.677e+05	
Df Model:	101					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	-1.11e+04	827.588	-13.414	0.000	-1.27e+04	-9478.875
sqft_living	92.5485	13.153	7.036	0.000	66.768	118.329
sqft_lot	0.2927	0.034	8.699	0.000	0.227	0.359
sqft_basement	24.3504	13.028	1.869	0.062	-1.185	49.886
yr_built	-645.8462	55.246	-11.690	0.000	-754.132	-537.560
lat	1.414e+05	5.54e+04	2.553	0.011	3.28e+04	2.5e+05
long	-1.861e+05	3.98e+04	-4.673	0.000	-2.64e+05	-1.08e+05
sqft_above	53.1626	13.180	4.034	0.000	27.329	78.996
yr_renovated	6.9696	1.200	5.809	0.000	4.618	9.321
sqft_living15	38.7259	2.565	15.098	0.000	33.699	43.753
sqft_lot15	-0.0845	0.054	-1.568	0.117	-0.190	0.021
bathrooms[0]	6819.1897	597.974	11.404	0.000	5647.115	7991.265
bathrooms[1]	6819.1897	597.974	11.404	0.000	5647.115	7991.265
bathrooms[2]	6819.1897	597.974	11.404	0.000	5647.115	7991.265
bathrooms[3]	6819.1897	597.974	11.404	0.000	5647.115	7991.265
floors[0]	-5276.1961	703.622	-7.499	0.000	-6655.348	-3897.044
floors[1]	-5276.1961	703.622	-7.499	0.000	-6655.348	-3897.044
floors[2]	-5276.1961	703.622	-7.499	0.000	-6655.348	-3897.044
floors[3]	-5276.1961	703.622	-7.499	0.000	-6655.348	-3897.044
waterfront[0]	1.824e+05	3149.220	57.905	0.000	1.76e+05	1.89e+05

waterfront[1]	1.824e+05	3149.220	57.905	0.000	1.76e+05	1.89e+05
waterfront[2]	1.824e+05	3149.220	57.905	0.000	1.76e+05	1.89e+05
waterfront[3]	1.824e+05	3149.220	57.905	0.000	1.76e+05	1.89e+05
condition[0]	1.374e+04	2884.312	4.762	0.000	8082.786	1.94e+04
condition[1]	1.374e+04	2884.312	4.762	0.000	8082.786	1.94e+04
condition[2]	1.374e+04	2884.312	4.762	0.000	8082.786	1.94e+04
grade[0]	1.729e+04	6704.427	2.579	0.010	4146.463	3.04e+04
grade[1]	1.729e+04	6704.427	2.579	0.010	4146.463	3.04e+04
grade[2]	1.729e+04	6704.427	2.579	0.010	4146.463	3.04e+04
zipcode[0]	-96.6712	18.290	-5.286	0.000	-132.520	-60.822
zipcode[1]	-96.6712	18.290	-5.286	0.000	-132.520	-60.822
zipcode[2]	-96.6712	18.290	-5.286	0.000	-132.520	-60.822
bed_2	1.102e+04	7955.804	1.385	0.166	-4573.723	2.66e+04
bed_3	1.912e+04	5269.881	3.628	0.000	8791.674	2.95e+04
bed_4	6408.2021	3608.259	1.776	0.076	-664.258	1.35e+04
bed_5	866.7682	4297.550	0.202	0.840	-7556.754	9290.290
bed_6	-1.265e+04	6417.575	-1.972	0.049	-2.52e+04	-74.548
bedrooms[0]	-2414.4030	1015.512	-2.378	0.017	-4404.884	-423.922
bedrooms[1]	-2414.4030	1015.512	-2.378	0.017	-4404.884	-423.922
bedrooms[2]	-2414.4030	1015.512	-2.378	0.017	-4404.884	-423.922
condi_2	3.659e+04	2.03e+04	1.800	0.072	-3248.890	7.64e+04
condi_3	1.269e+04	9132.466	1.390	0.165	-5206.924	3.06e+04
condi_4	-5780.6609	2222.139	-2.601	0.009	-1.01e+04	-1425.101
condi_5	-4948.6116	8659.385	-0.571	0.568	-2.19e+04	1.2e+04
grd_4	-1.475e+05	1.23e+05	-1.199	0.231	-3.89e+05	9.36e+04
grd_5	-2.439e+05	1.01e+05	-2.425	0.015	-4.41e+05	-4.68e+04
grd_6	-2.99e+05	8.04e+04	-3.718	0.000	-4.57e+05	-1.41e+05
grd_7	-3.424e+05	6.06e+04	-5.649	0.000	-4.61e+05	-2.24e+05
grd_8	-3.592e+05	4.11e+04	-8.735	0.000	-4.4e+05	-2.79e+05
grd_9	-3.234e+05	2.27e+04	-14.231	0.000	-3.68e+05	-2.79e+05
grd_10	-2.499e+05	1.24e+04	-20.138	0.000	-2.74e+05	-2.26e+05
grd_11	-9.688e+04	2.48e+04	-3.901	0.000	-1.46e+05	-4.82e+04
grd_12	1.869e+05	4.48e+04	4.174	0.000	9.91e+04	2.75e+05
grd_13	7.532e+05	7.59e+04	9.923	0.000	6.04e+05	9.02e+05
zip_c_98002	2.355e+04	1.26e+04	1.862	0.063	-1242.622	4.83e+04
zip_c_98003	-9671.8789	1.12e+04	-0.860	0.390	-3.17e+04	1.24e+04
zip_c_98004	7.013e+05	2.05e+04	34.175	0.000	6.61e+05	7.42e+05
zip_c_98005	2.666e+05	2.19e+04	12.181	0.000	2.24e+05	3.09e+05
zip_c_98006	2.429e+05	1.79e+04	13.538	0.000	2.08e+05	2.78e+05
zip_c_98007	2.203e+05	2.26e+04	9.731	0.000	1.76e+05	2.65e+05

zip_c_98008	2.5e+05	2.15e+04	11.646	0.000	2.08e+05	2.92e+05
zip_c_98010	1.126e+05	1.95e+04	5.780	0.000	7.44e+04	1.51e+05
zip_c_98011	7.668e+04	2.78e+04	2.758	0.006	2.22e+04	1.31e+05
zip_c_98014	1.226e+05	3.09e+04	3.966	0.000	6.2e+04	1.83e+05
zip_c_98019	9.223e+04	3.03e+04	3.042	0.002	3.28e+04	1.52e+05
zip_c_98022	8.253e+04	1.75e+04	4.707	0.000	4.82e+04	1.17e+05
zip_c_98023	-4.345e+04	1e+04	-4.323	0.000	-6.31e+04	-2.37e+04
zip_c_98024	1.843e+05	2.75e+04	6.703	0.000	1.3e+05	2.38e+05
zip_c_98027	1.758e+05	1.87e+04	9.409	0.000	1.39e+05	2.12e+05
zip_c_98028	7.822e+04	2.67e+04	2.934	0.003	2.6e+04	1.3e+05
zip_c_98029	2.362e+05	2.13e+04	11.069	0.000	1.94e+05	2.78e+05
zip_c_98030	1.705e+04	1.27e+04	1.347	0.178	-7759.708	4.19e+04
zip_c_98031	2.181e+04	1.31e+04	1.665	0.096	-3871.639	4.75e+04
zip_c_98032	638.5251	1.48e+04	0.043	0.966	-2.85e+04	2.97e+04
zip_c_98033	3.278e+05	2.29e+04	14.292	0.000	2.83e+05	3.73e+05
zip_c_98034	1.651e+05	2.45e+04	6.739	0.000	1.17e+05	2.13e+05
zip_c_98038	7.731e+04	1.5e+04	5.144	0.000	4.79e+04	1.07e+05
zip_c_98039	1.061e+06	2.9e+04	36.561	0.000	1e+06	1.12e+06
zip_c_98040	4.739e+05	1.78e+04	26.551	0.000	4.39e+05	5.09e+05
zip_c_98042	3.576e+04	1.28e+04	2.793	0.005	1.07e+04	6.09e+04
zip_c_98045	1.887e+05	2.71e+04	6.953	0.000	1.35e+05	2.42e+05
zip_c_98052	2.176e+05	2.35e+04	9.245	0.000	1.72e+05	2.64e+05
zip_c_98053	2.131e+05	2.56e+04	8.312	0.000	1.63e+05	2.63e+05
zip_c_98055	4.886e+04	1.44e+04	3.382	0.001	2.05e+04	7.72e+04
zip_c_98056	9.577e+04	1.56e+04	6.144	0.000	6.52e+04	1.26e+05
zip_c_98058	4.578e+04	1.42e+04	3.227	0.001	1.8e+04	7.36e+04
zip_c_98059	9.06e+04	1.57e+04	5.765	0.000	5.98e+04	1.21e+05
zip_c_98065	1.62e+05	2.53e+04	6.408	0.000	1.12e+05	2.12e+05
zip_c_98070	8584.9938	1.64e+04	0.523	0.601	-2.36e+04	4.08e+04
zip_c_98072	1.294e+05	2.73e+04	4.741	0.000	7.59e+04	1.83e+05
zip_c_98074	1.815e+05	2.3e+04	7.903	0.000	1.36e+05	2.26e+05
zip_c_98075	1.905e+05	2.24e+04	8.495	0.000	1.47e+05	2.34e+05
zip_c_98077	9.471e+04	2.88e+04	3.286	0.001	3.82e+04	1.51e+05
zip_c_98092	1.762e+04	1.37e+04	1.283	0.200	-9302.858	4.46e+04
zip_c_98102	4.238e+05	2.26e+04	18.767	0.000	3.8e+05	4.68e+05
zip_c_98103	3.053e+05	2.04e+04	14.986	0.000	2.65e+05	3.45e+05
zip_c_98105	4.448e+05	2.16e+04	20.587	0.000	4.02e+05	4.87e+05
zip_c_98106	1.132e+05	1.47e+04	7.692	0.000	8.44e+04	1.42e+05
zip_c_98107	3.039e+05	2.08e+04	14.589	0.000	2.63e+05	3.45e+05
zip_c_98108	1.075e+05	1.71e+04	6.284	0.000	7.4e+04	1.41e+05

zip_c_98109	4.667e+05	2.2e+04	21.177	0.000	4.24e+05	5.1e+05
zip_c_98112	5.745e+05	1.97e+04	29.138	0.000	5.36e+05	6.13e+05
zip_c_98115	3.118e+05	2.1e+04	14.851	0.000	2.71e+05	3.53e+05
zip_c_98116	2.898e+05	1.6e+04	18.125	0.000	2.58e+05	3.21e+05
zip_c_98117	2.811e+05	2.05e+04	13.732	0.000	2.41e+05	3.21e+05
zip_c_98118	1.695e+05	1.53e+04	11.068	0.000	1.39e+05	1.99e+05
zip_c_98119	4.552e+05	2.03e+04	22.399	0.000	4.15e+05	4.95e+05
zip_c_98122	3.344e+05	1.86e+04	17.986	0.000	2.98e+05	3.71e+05
zip_c_98125	1.836e+05	2.27e+04	8.097	0.000	1.39e+05	2.28e+05
zip_c_98126	1.899e+05	1.47e+04	12.956	0.000	1.61e+05	2.19e+05
zip_c_98133	1.273e+05	2.29e+04	5.563	0.000	8.24e+04	1.72e+05
zip_c_98136	2.596e+05	1.49e+04	17.391	0.000	2.3e+05	2.89e+05
zip_c_98144	2.884e+05	1.72e+04	16.734	0.000	2.55e+05	3.22e+05
zip_c_98146	1.176e+05	1.37e+04	8.580	0.000	9.08e+04	1.45e+05
zip_c_98148	8.107e+04	2.09e+04	3.886	0.000	4.02e+04	1.22e+05
zip_c_98155	1.215e+05	2.41e+04	5.037	0.000	7.42e+04	1.69e+05
zip_c_98166	9.899e+04	1.33e+04	7.469	0.000	7.3e+04	1.25e+05
zip_c_98168	7.003e+04	1.46e+04	4.806	0.000	4.15e+04	9.86e+04
zip_c_98177	2.2e+05	2.33e+04	9.446	0.000	1.74e+05	2.66e+05
zip_c_98178	7.61e+04	1.64e+04	4.644	0.000	4.4e+04	1.08e+05
zip_c_98188	7.111e+04	1.7e+04	4.173	0.000	3.77e+04	1.05e+05
zip_c_98198	6.279e+04	1.38e+04	4.566	0.000	3.58e+04	8.98e+04
zip_c_98199	3.825e+05	1.83e+04	20.856	0.000	3.47e+05	4.18e+05


Omnibus:	12606.338	Durbin-Watson:	1.978
Prob(Omnibus):	0.000	Jarque-Bera (JB):	398788.018
Skew:	2.296	Prob(JB):	0.00
Kurtosis:	23.662	Cond. No.	7.20e+18

Notes:

- [1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [2] The smallest eigenvalue is 1.21e-23. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

In [325]:

```
# Q-Q plot and Regression plot
model=smf.ols(formula,data=d1).fit()
fig = sm.graphics.qqplot(model.resid, line='45',fit=True);
#fig.savefig('Basic_model_qqplot')
sm.graphics.plot_regress_exog(model, 'sqft_living', fig=plt.figure(figsize=(12,8)));
```



Conclusion:

The base model shows that sqft_living(independent variable) has high correlation with price(dependent variable). In Basic linear model high R-squared value (0.82) and low P-value(0) shows there is a relationship between variables. Q-Q data plot is not normally distributed and scatter plots are very homoscedastic.

Iteration1

Log transformation

In [326]:

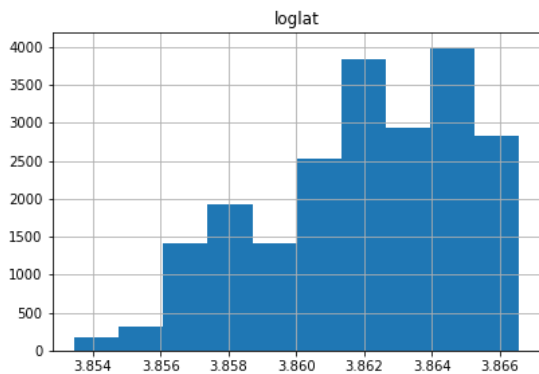
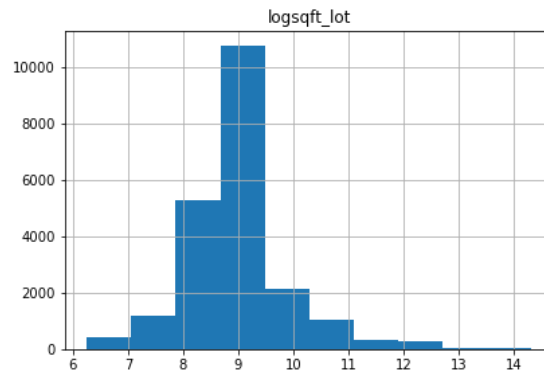
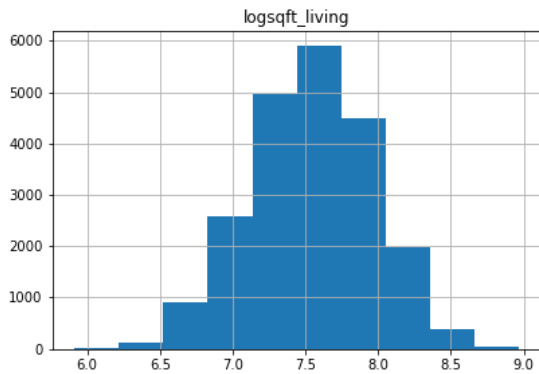


```
n.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 21597 entries, 0 to 21596
Data columns (total 13 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   bedrooms        21363 non-null  float64
 1   bathrooms        21363 non-null  float64
 2   floors           21363 non-null  float64
 3   waterfront       21363 non-null  float64
 4   condition        21363 non-null  float64
 5   grade            21363 non-null  float64
 6   zipcode          21363 non-null  float64
 7   price            21597 non-null  float64
 8   sqft_living      21597 non-null  int64
 9   sqft_lot         21597 non-null  int64
10  yr_built         21597 non-null  int64
11  lat              21597 non-null  float64
12  long             21597 non-null  float64
dtypes: float64(10), int64(3)
memory usage: 2.9 MB
```

In [346]:

```
data_log = pd.DataFrame([])
data_log['logsqft_living'] = np.log(d1['sqft_living'])
data_log['logsqft_lot'] = np.log(d1['sqft_lot'])
#data_log['logsqft_basement'] = np.log(d1['sqft_basement'])
data_log['loglat'] = np.log(d1['lat'])
#data_log['loglong'] = np.log(d1['long'])
data_log.hist(figsize = [15,10 ]);
```



In [347]:

```
pp = d1['price']
logliving = np.log(d1['sqft_living'])
loglot = np.log(d1['sqft_lot'])
loglat = np.log(d1['lat'])
#loglong = np.log(d1['long'])

scaled_pp = (pp-min(pp))/(max(pp)-min(pp))
scaled_lot = (loglot-np.mean(loglot))/np.sqrt(np.var(loglot))
scaled_lat = (loglat-np.mean(loglat))/(max(loglat)-min(loglat))
scaled_living=(logliving-np.mean(logliving))/np.sqrt(np.var(logliving))

data_fin = pd.DataFrame([])
data_fin['pp'] = scaled_pp
data_fin['lot'] = scaled_lot
data_fin['lat'] = scaled_lat
data_fin['living'] = scaled_living
```

In [348]:



```
data_fin
```

Out[348]:

	pp	lot	lat	living
0	0.032616	-0.382321	-0.077514	-1.126726
1	0.104261	-0.105936	0.258885	0.746830
2	0.023119	0.253334	0.285919	-2.154206
3	0.119220	-0.518395	-0.062088	0.094637
4	0.097915	0.015972	0.091993	-0.276397
...
21592	0.063917	-2.173235	0.224160	-0.501510
21593	0.072983	-0.350656	-0.078317	0.490108
21594	0.073459	-1.976166	0.056069	-1.477448
21595	0.072983	-1.341154	-0.040081	-0.393833
21596	0.055984	-2.228739	0.055588	-1.477448

21363 rows × 4 columns

In [349]:



```
outcome = 'pp'
x_cols = data_fin.drop('pp',axis=1)
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=data_fin).fit()
model.summary()
```

Out[349]:

OLS Regression Results

Dep. Variable:	pp		R-squared:	0.469		
Model:	OLS		Adj. R-squared:	0.469		
Method:	Least Squares		F-statistic:	6292.		
Date:	Sun, 12 Feb 2023		Prob (F-statistic):	0.00		
Time:	22:23:01		Log-Likelihood:	31874.		
No. Observations:	21363		AIC:	-6.374e+04		
Df Residuals:	21359		BIC:	-6.371e+04		
Df Model:	3					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	0.1023	0.000	274.582	0.000	0.102	0.103
lot	0.0003	0.000	0.872	0.383	-0.000	0.001
lat	0.1044	0.002	61.560	0.000	0.101	0.108
living	0.0447	0.000	113.442	0.000	0.044	0.045
Omnibus:	15259.012	Durbin-Watson:	1.985			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	396659.314			
Skew:	3.137	Prob(JB):	0.00			
Kurtosis:	23.156	Cond. No.	5.23			

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

In [350]:



```
r_lot = data_fin.drop('lot',axis=1)
r_lot
```

Out[350]:

	pp	lat	living
0	0.032616	-0.077514	-1.126726
1	0.104261	0.258885	0.746830
2	0.023119	0.285919	-2.154206
3	0.119220	-0.062088	0.094637
4	0.097915	0.091993	-0.276397
...
21592	0.063917	0.224160	-0.501510
21593	0.072983	-0.078317	0.490108
21594	0.073459	0.056069	-1.477448
21595	0.072983	-0.040081	-0.393833
21596	0.055984	0.055588	-1.477448

21363 rows × 3 columns

In [351]:

```
outcome = 'pp'
x_cols = r_lot.drop('pp',axis=1)
predictors = '+'.join(x_cols)
formula = outcome + '~' + predictors
model = ols(formula=formula, data=r_lot).fit()
model.summary()
```


Out[351]:

OLS Regression Results

Dep. Variable:	pp	R-squared:	0.469
Model:	OLS	Adj. R-squared:	0.469
Method:	Least Squares	F-statistic:	9437.
Date:	Sun, 12 Feb 2023	Prob (F-statistic):	0.00
Time:	22:26:31	Log-Likelihood:	31873.
No. Observations:	21363	AIC:	-6.374e+04
Df Residuals:	21360	BIC:	-6.372e+04
Df Model:	2		
Covariance Type:	nonrobust		

In [353]:

```
# Q-Q plot and Regression plot
model=smf.ols(formula,data=r_lot).fit()
fig = sm.graphics.qqplot(model.resid, line='45',fit=True);
#fig.savefig('Basic_model_qqplot')
sm.graphics.plot_regress_exog(model, 'living', fig=plt.figure(figsize=(12,8)));
```



Conclusion

The Iteration 1 shows that sqft_living(independent variable) has high correlation with price(dependent variable). In this model high R-squared value (0.40) and low P-value(0) shows there is a relationship between variables. Q-Q data plot is not normally distributed and scatter plots are very heteroscedastic.

In [354]:

```
data_pred = data_fin.iloc[:,0:15]
data_pred
```

Out[354]:

	pp	lot	lat	living
0	0.032616	-0.382321	-0.077514	-1.126726
1	0.104261	-0.105936	0.258885	0.746830
2	0.023119	0.253334	0.285919	-2.154206
3	0.119220	-0.518395	-0.062088	0.094637
4	0.097915	0.015972	0.091993	-0.276397
...
21592	0.063917	-2.173235	0.224160	-0.501510
21593	0.072983	-0.350656	-0.078317	0.490108
21594	0.073459	-1.976166	0.056069	-1.477448
21595	0.072983	-1.341154	-0.040081	-0.393833

In [355]:

```
data_pred.corr()
```

Out[355]:

	pp	lot	lat	living
pp	1.000000	0.146261	0.331163	0.610355
lot	0.146261	1.000000	-0.152340	0.316015
lat	0.331163	-0.152340	1.000000	0.033681
living	0.610355	0.316015	0.033681	1.000000

In [356]:

```
abs(data_pred.corr()) > 0.5
```

Out[356]:

	pp	lot	lat	living
pp	True	False	False	True
lot	False	True	False	False
lat	False	False	True	False
living	True	False	False	True

In [357]:

```
plt.figure(figsize=(15,10))
sns.heatmap(data_fin.corr().round(2),annot=True)
```

Out[357]:

<AxesSubplot:>



In []:

Conclusion:

The price and sqft living is 61% correlated and showing normal distribution.
the r2 and adjust r2 is 45%

Iteration 2-Model evaluation

Train-test split

In [335]:

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import FunctionTransformer, OneHotEncoder
from sklearn.linear_model import LinearRegression
```

In [336]:



```
X = cont_data.drop("price", axis=1)
y = cont_data["price"]
```

In [337]:



```
X_train, X_test, y_train, y_test = train_test_split(X,y,test_size=0.3,random_state=101)
```

In [338]:



```
lm=LinearRegression()
lm.fit(X_train,y_train)
```

Out[338]:

```
LinearRegression()
```

In [339]:



```
rootMeanSquareError = np.sqrt(meanSquareError).round(3)
rootMeanSquareError
```

Out[339]:

```
228036.996
```

In [340]:



```
y_predict= lm.predict(X_train)
```

In [341]:



```
y_predict
```

Out[341]:

```
array([765179.34565939, 288045.7561236 , 694204.53102167, ...,
       412681.85528234, 357586.28744134, 294053.1575178 ])
```

In [342]:



```
#model evaluation
print('R^2:',metrics.r2_score(y_train, y_predict))
print('Adjusted R^2:',1 - (1-metrics.r2_score(y_train, y_predict))*(len(y_train)-1)/(len(y_
print('MAE:',metrics.mean_absolute_error(y_train, y_predict))
print('MSE:',metrics.mean_squared_error(y_train, y_predict))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_train, y_predict)))
```


```
R^2: 0.6003850414249247
Adjusted R^2: 0.6001205008724455
MAE: 148100.68136135113
MSE: 54323037409.31655
RMSE: 233073.030205806
```

In [343]:



```
import matplotlib.pyplot as plt

plt.scatter(y_train,y_predict)
plt.title("Prices vs Predicted prices")
plt.show()
```



In [344]:

```
#Linear regression for ols

y_test_pred=lm.predict(X_test)


acc_linreg=metrics.r2_score(y_test,y_test_pred)
print('R^2:',acc_linreg)
print('Adjusted R^2:', 1- (1-metrics.r2_score(y_test,y_test_pred))*(len(y_test)-1)/(len(y_t
print('MAE:',metrics.mean_absolute_error(y_test, y_test_pred))
print('MSE:',metrics.mean_squared_error(y_test, y_test_pred))
print('RMSE:',np.sqrt(metrics.mean_squared_error(y_test, y_test_pred)))
```

```
R^2: 0.6166708534334627
Adjusted R^2: 0.6160782902141606
MAE: 143292.99656191687
MSE: 50835736790.27583
RMSE: 225467.81763763056
```

In [345]:

```
import matplotlib.pyplot as plt

plt.scatter(y_test,y_test_pred)
plt.title("Prices vs Predicted prices")
plt.show()
```



Conclusion:

This model shows that MSE value is higher for x-test than for x-train. R-squared value (0.60) for X-train and R-squared value (0.61) for X-test. The scatter plot is showing non linear effect.

In []:

